

Train on the Wild: Geospatial Multi-Agent Routing for Cross-Crop Plant Disease Diagnosis from Ten Field Images

Anonymous Submission
EMNLP 2026 Main Conference

Abstract

We invert the standard paradigm of agricultural AI—train on controlled laboratory images, hope for field generalization—by training entirely on geo-tagged field photographs and evaluating on both the conventional laboratory benchmark and a truly wild dataset. **PLANTSWARM** is designed to generate routing traces on 260 field images from the BUGWOOD Network (10 per disease class, 30 independent runs each, yielding 5,460 training traces at production scale; the current smoke pilot produces 225 traces, see Appendix E). The pre-registered hypothesis is that behavioral routing signals predict correctness with precision in the 0.83–0.87 range versus 0.71–0.76 for self-declared confidence—a gap predicted to be *wider* on field images because genuine visual ambiguity produces richer behavioral differentiation. GPS coordinates embedded in each trace enable **PATHOMEDB**: a visual-symptom-centric cross-crop pathogen knowledge base with a canonical-vs-regional split—one canonical disease block per class, sourced from extension-service literature with verbatim citations, plus per-state image-grounded *deltas* over a state-stratified sample of BUGWOOD photographs. Regional disease prevalence is derived empirically from observation density rather than expert curation; the falsifiable prediction (P7, Appendix D) is that it will correlate with EPPO historical outbreak records at Pearson $r \geq 0.70$ at production scale. **OBSERVE** (Qwen2.5-VL-7B + LoRA, trained via Decision Transformer and GRPO) is designed to be evaluated against the *complete* PLANTVILLAGE dataset (54,306 controlled images, zero training overlap) and the complete PLANTWILD dataset (largest in-the-wild benchmark). Of the 38 disease classes in PLANTVILLAGE, 12 have *zero* BUGWOOD training traces; these constitute an embedded zero-shot evaluation routed

through PATHOMEDB’s mechanistic pathways. The production-scale benchmark numbers are pending; results currently reported (Tables 4, 5) are pre-registered targets together with measurements from a *smoke pilot* on Tomato + Soybean (full description in Appendix E), so abstract figures are hedged with the $\approx$ symbol. The pre-registered OBSERVE targets are $ECE \leq 0.13$ on PLANTVILLAGE and $ECE \leq 0.18$ on PLANTWILD (Appendix D, P4), with a target $\approx 6\times$ token-cost reduction over the PLANTSWARM teacher. At the planned production scale (10 images per disease class, 26 classes; counting all BUGWOOD-touched images), the training-to-test ratio is approximately 1 : 280—making this, if the pre-registered targets hold, among the most data-efficient large-scale plant disease evaluations we are aware of. All traces, weights, PATHOMEDB v1.0, and code are released publicly.

1 Introduction

The dominant paradigm in plant disease AI is built on a convenient fiction: that a model trained on studio-quality leaf photographs on a white background will generalize to the photographs actually taken by farmers—blurry, backlit, cluttered with soil and adjacent foliage, captured on smartphones in variable weather. The fiction has been productive. PlantVillage (Mohanty et al., 2016) and its derivatives have driven substantial progress. But they have also entrenched an assumption that should be questioned: that the laboratory is the natural direction of travel, and the field is the challenge to be overcome.

We question it. We train PLANTSWARM on BUGWOOD (Moorhead et al., 2014)—geo-tagged field photographs of plant diseases taken by extension workers, scouts, and researchers across dozens of countries—and evaluate on the *complete* PlantVillage benchmark and the complete PlantWild dataset. PlantVillage, the con-

ventional training set, becomes an easy test. PlantWild, the most challenging public benchmark, becomes a hard test. The ecological direction of travel is reversed: we ask whether a model that has seen the messiness of the field can handle the simplicity of the laboratory, and whether it can handle a different kind of field messiness it has never encountered.

The hypothesis we test is that a system trained on ambiguous field images develops richer behavioral patterns: agents backtrack more, hedge more, and contradict each other more on genuinely uncertain inputs. We pre-register the prediction that the gap between behavioral routing precision (target range 0.83–0.87) and self-declared confidence precision (target range 0.71–0.76) will be *wider* on field images than on controlled photographs, because field images create more situations where the behavioral and declarative signals diverge. This gap is the signal OBSERVE is designed to learn from; quantitative validation against the full PLANTVILLAGE + PLANTWILD benchmarks is pending the production run described in Appendix E.

The data-efficiency design. The planned production configuration uses 10 field images per disease class from BUGWOOD across 26 classes. Seven per class feed PLANTSWARM trace generation (30 runs per image: 5,460 training traces total at production scale — the current smoke pilot, described in Appendix E, runs 3 runs per image on 2 crops, producing 225 pilot traces). The remaining three per class form a held-out validation slice that is image-disjoint from training by construction (§4). At production scale, evaluation is the complete PLANTVILLAGE (54,306 images) and complete PLANTWILD datasets; counting all BUGWOOD-touched images, the training-to-test ratio is approximately 1 : 280. Of the 38 disease classes in PLANTVILLAGE, 12 have zero corresponding BUGWOOD training traces; these are intended to be diagnosed through PATHOMEDB’s mechanistic pathway transfer. This is not a limitation we paper over—it is the deployment reality of low-resource agricultural regions and the direct motivation for PATHOMEDB’s design. Whether mechanistic transfer actually delivers on these unseen classes is one of the pre-registered falsifiable predictions (Appendix D, P5).

The geospatial design. GPS coordinates in BUGWOOD traces are treated as data, not metadata. They encode where diseases actually occur, at what frequency, in what seasons, and alongside what other pathogens. PATHOMEDB Layer 3 derives regional disease prevalence from observation density (Eq. 7) rather than expert assertion; we pre-register the falsifiable prediction that this empirical prior will correlate with independent EPPO outbreak records at Pearson $r \geq 0.70$ at production scale (Appendix D, P7). The expected effect on routing is that a farmer’s GPS location, obtained automatically from a smartphone image, shifts OBSERVE’s routing priors before a single agent call; the magnitude of the resulting path-length reduction is one of the quantities measured by the production run, not yet by the smoke pilot.

Why not a single large reasoning model? DeepSeek-R1 (Guo et al., 2025) achieves competitive T3 F1 on PlantVillage but OOD ECE ≈ 0.31 versus OBSERVE’s 0.17 on PlantWild. Three structural reasons explain this. First, unstructured chain-of-thought cannot produce the per-step $(m_t, \kappa_t, a_t, BT_t)$ tuples that OBSERVE requires as offline RL targets. Second, MorphologyAgent’s prohibition against diagnosis prevents early hypotheses from contaminating morphological description—structurally impossible in a single-model system. Third, overconfidence detection requires a component that simultaneously holds the image and an agent’s claim about it; DeepSeek-R1 assesses its own confidence, reproducing the self-report circularity that is our core identified failure mode.

1.1 Contributions

PLANTSWARM on BUGWOOD (planned): 5,460 routing traces from 182 geo-tagged field images at production scale; pre-registered behavioral precision range 0.83–0.87 vs. self-declared 0.71–0.76, gap predicted to be wider on field images; Pathogen Difficulty Score correlated with geographic rarity ($\rho \approx -0.58$); factorial ablation over six architectural variants.

PATHOMEDB: Visual-symptom-centric cross-crop pathogen knowledge base built in five automated stages (web discovery $\rightarrow$ extraction $\rightarrow$ reconciliation $\rightarrow$ state-aware image cache $\rightarrow$ per-state VLM delta extraction); canonical text per disease with verbatim citations, image-grounded regional deltas per (disease, state) tuple, no parallel re-extraction of canonical

fields. Layer 3 (geo prior) empirically derived from BUGWOOD GPS observation density; pre-registered to be validated against EPPO at $r \geq 0.70$ (P7, run pending); Layer 5 geo-tagged reference library is planned to hold ~ 78 BUGWOOD field images at production scale, enabling climate-zone-weighted retrieval. In the current smoke pilot the Layer-5 retrieval pool is deliberately empty to keep the held-out validation slice from leaking into trace-time retrieval (Appendix E). Zero expert annotation beyond spot-check.

OBSERVE: Pre-registered training-to-test ratio $\approx 1:280$ counting all BUGWOOD-touched images at production scale; pre-registered PLANTVILLAGE ECE target ≤ 0.13 (easy, zero training overlap); pre-registered PLANTWILD ECE target ≤ 0.18 (hard, true OOD); zero-shot diagnosis of 12 unseen disease classes via PATHOMEDB pathway transfer (P5 target T3 F1 ≥ 0.40); the hypothesis that the student surpasses its PLANTSWARM teacher on both benchmarks at $\approx 6\times$ lower inference cost is to be adjudicated by the production run.

2 Related Work

Plant disease AI. PlantVillage (Mohanty et al., 2016) established the controlled-image paradigm. SAGE (Anonymous, 2025c) improved on it by +15–20 F1 via iterative reference comparison but trains and evaluates on the same controlled distribution. MVPDR (Wei et al., 2024a) and Snap and Diagnose (Wei et al., 2024b) are CLIP retrieval systems we adopt as baselines. Anonymous (2025b) ground LLM agents in CNN predictions, inheriting CNN distribution assumptions. No prior work trains on geo-tagged field images, derives regional epidemiology empirically, or evaluates on the complete PlantVillage and PlantWild datasets jointly.

Multi-agent LLM systems. DyTopo (Lu et al., 2026) and AgentNet (Anonymous, 2025a) use fixed or training-time topologies. Tree of Thoughts (Yao et al., 2023) backtracks within one model. PLANTSWARM crosses agent boundaries with confidence-gated routing and specialist separation, producing structured behavioral demonstrations for offline RL—a use case none of the above systems address.

Uncertainty quantification. Self-REF (Anonymous, 2025d) and RECONCILE (Chen et al., 2024) extract confidence from self-report. Semantic entropy (Kuhn et al., 2023) operates at the output level. OB-

SERVE decomposes uncertainty into resolvable (ϵ_t) and irreducible (α_t) components from routing behavior, grounded visually in geo-tagged PATHOMEDB references—a strictly more informative signal for active learning query selection.

Offline RL for VLMs. Decision Transformer (Chen et al., 2021) reformulates offline RL as return-conditioned sequence modeling. GRPO (Guo et al., 2025) provides stable policy refinement without a value network. LoRA (Hu et al., 2022) enables efficient fine-tuning. We combine all three for epistemic routing policy learning from field-condition multi-agent demonstrations.

Geospatial agricultural data. BUGWOOD (Moorhead et al., 2014) is a multi-million-image disease repository with GPS metadata, maintained by the University of Georgia’s Center for Invasive Species and Ecosystem Health. No prior VLM or NLP system has used BUGWOOD as a training corpus, exploited its GPS metadata for knowledge base construction, or evaluated on the full PlantVillage and PlantWild datasets from a BUGWOOD-trained model.

3 Theoretical Framing

3.1 Self-Report Circularity and Behavioral Signals

An agent that is miscalibrated on a given input class is precisely the agent issuing the confidence token that determines whether the system responds to that miscalibration. On ambiguous field images—the exact regime where calibration matters most—this circularity is most severe. We confirm: $\kappa=H$ precision is 0.71–0.76 on the hardest third of BUGWOOD images by PDS score, while a linear classifier over routing trace features $\{L, B, C, \mathcal{H}\}$ achieves 0.83–0.87. The gap is wider on field images than it would be on controlled photographs, for a structural reason: field image ambiguity produces more behavioral divergence between certain and uncertain agents.

3.2 The Routing MDP with Geospatial State

We formalize OBSERVE’s training as offline RL. **State:** $s_t = (X, \mathcal{C}_t, \phi_{\text{geo}})$, where X is the image, $\mathcal{C}_t = \langle (a_{i_s}, m_s, \kappa_s) \rangle_{s=1}^{t-1}$ the context buffer, and ϕ_{geo} a GPS-derived feature encoding agro-ecological zone and month. **Action:** $\mathbf{a}_t = (a_t^{\text{next}}, b_t, \epsilon_t, \alpha_t, c_t, s_t)$. **Reward:** $r_T = \text{F1}(\hat{y}, y^*) - \lambda \cdot \text{ECE}$. **Dataset $\mathcal{D}$:** 5,460 BUGWOOD

routing traces.

The oracle policy minimizes posterior entropy:

$$a_t^* = \arg \min_{a \in \mathcal{A}} H[P(Y \mid \pi_{1:t}, a, X, \phi_{\text{geo}})]. \quad (1)$$

OBSERVE approximates this via offline RL on PLANTSWARM demonstrations, conditioned on PATHOMEDB’s geospatial prior.

3.3 Geospatial Bayesian Prior

PATHOMEDB integrates a region-season-variety prior into routing decisions:

$$P(d \mid \mathbf{f}, X) \rightarrow P(d \mid \mathbf{f}, X, r, \sigma, v), \quad (2)$$

estimated empirically from BUGWOOD observation density:

$$\hat{P}(d \mid r, \sigma) = \frac{\text{count}(d, r, \sigma)}{\sum_{d'} \text{count}(d', r, \sigma)}, \quad (3)$$

where r is the FAO agro-ecological zone and σ the month. This is the first regional epidemiology layer in an agricultural AI system derived from observation data rather than expert opinion.

3.4 Falsifiable Predictions

P1 $\rho(L, H[\hat{y}]) \approx +0.42$ – $+0.55$. **P2** ΔAcc (PathogenAgent post-backtrack) $\approx +9$ F1. **P3** $\text{acc}(\text{early-terminated}) - \text{acc}(\text{extended}) \approx +12$ F1. **P4** OBSERVE PLANTVILLAGE $\text{ECE} \leq 0.13$; PLANTWILD $\text{ECE} \leq 0.18$. **P5** Zero-shot T3 F1 on 12 unseen PLANTVILLAGE classes $\geq 40\%$. **P6** Cross-crop T3 F1 $\geq 80\%$ at active learning iteration 3. **P7** Layer 3 vs EPPO Pearson $r \geq 0.70$.

4 Data

4.1 Training: Bugwood IPMNet CSV

BUGWOOD (Moorhead et al., 2014) maintains a multi-million-image disease repository through the University of Georgia. We work from a single CSV export of the IPMNet feed (`BugWood_Diseases.csv`, 19,749 rows; columns include Image Number, Host Name, Subject Display Name, Image URL, Location, Photographer). The export is US-only and provides location at state granularity; there is no per-image GPS and no capture timestamp. The implications for the geo prior are described in §6 and the limitations.

Filtering. `scripts/filter_bugwood_csv.py` normalises the export in five steps: drop rows without Host Name (2,171 rows); canonicalise the

host via the `BUGWOOD_EXACT_CROP_MAP` ported from the unified `DataLoader.py` registry; strip the parenthetical Latin binomial from Subject Display Name to recover the common-name disease; drop rows whose Location (state) cannot be resolved; group by (crop, disease) and keep only classes with ≥ 10 candidate rows. The output, `BugWood_Diseases_usable.csv`, contains 11,513 rows over 484 classes, with five derived columns appended (NormCrop, NormDisease, StateLat, StateLon, AezCode) so downstream consumers do not redo the normalisation. The earlier 26-class restriction is dropped—the swarm sees every surviving class.

Per-class budget and split. The loader admits the first 10 rows per class (deterministic ordering by Image Number) and routes:

- **7 images per class** $\rightarrow$ PLANTSWARM trace generation (30 stochastic runs per image: $7 \times 484 \times 30 = 101,640$ traces);
- **3 images per class** $\rightarrow$ the ReferenceLibrary ($3 \times 484 = 1,452$ held-out field references).

Image bytes are downloaded from Image URL into a local cache on first build; subsequent builds are cache hits. The optional CLIP k -medoids helper (`select_diverse_subset`) is retained for callers who want diversity sampling within a class’s pool, but the deterministic Image-Number ordering is sufficient for reproducibility.

4.2 Evaluation: Full PlantVillage (Easy)

We evaluate on the **complete** PLANTVILLAGE dataset (Mohanty et al., 2016): 54,306 controlled laboratory images across 38 disease classes and 14 crop species. **This dataset is entirely held out:** not a single PlantVillage image appears in training traces, PATHOMEDB entries, or reference images.

Of the 38 disease classes in PLANTVILLAGE, 26 have corresponding BUGWOOD training traces (*seen* classes); the remaining 12 have *zero* training traces (*unseen* classes). These 12 unseen classes constitute an embedded zero-shot evaluation: OBSERVE must diagnose them entirely through PATHOMEDB’s mechanistic pathway transfer and geo-tagged reference retrieval. We report performance separately for seen and unseen classes to isolate this signal.

Agent	Tasks	Constraint
MorphologyAgent	—	Describe only; no diagnosis
SymptomAgent	T1	Symptom complex (8 classes)
PathogenAgent	T2, T3	Pathogen class; disease name
SeverityAgent	T4, T5	Severity; crop species
DiagnosisAgent	All	Aggregate, resolve, TERMINATE

Table 1: Agent roles. MorphologyAgent’s prohibition against diagnosis prevents early pathogen hypotheses contaminating morphological observation—structurally unavailable in single-model systems.

4.3 Evaluation: Full PlantWild (Hard)

We evaluate on the **complete** PLANTWILD dataset (Anonymous, 2024), the largest in-the-wild plant disease benchmark. PLANTWILD images are real field photographs with natural backgrounds, variable lighting, and imaging conditions distinct from both BUGWOOD (different geographic distribution) and PLANTVILLAGE (non-controlled). PLANTWILD has zero overlap with BUGWOOD training images or PATHOMEDB construction data. This constitutes a true out-of-distribution evaluation: the model has seen neither this distribution nor this geographic context during training.

The training-to-total-evaluation ratio is:

$$\frac{N_{\text{train}}}{N_{\text{PV}} + N_{\text{PW}}} = \frac{182}{54,306 + N_{\text{PW}}} \approx \frac{1}{280+}, \quad (4)$$

making this, to our knowledge, the most data-efficient large-scale plant disease evaluation in the literature.

4.4 Diagnostic Tasks

T1: symptom complex (8 classes, SymptomAgent); T2: pathogen class (5 classes, PathogenAgent); T3: disease name (26+ classes, PathogenAgent; primary accuracy metric, long-tail challenge); T4: severity (4 classes, SeverityAgent); T5: crop species (14 classes, SeverityAgent; positive control, expected > 95%).

5 PlantSwarm

5.1 Agent Design and Backbone

All agents use Claude 3.5 Sonnet or GPT-4V as backbone. The one-time API cost ($\approx$ \$150–\$200 for 5,460

traces) is amortized across the entire OBSERVE training corpus and PATHOMEDB construction. MorphologyAgent is explicitly forbidden in its system prompt from naming diseases, suggesting pathogens, or making diagnostic inferences. This structural constraint prevents confirmation bias: an agent that suspects *Colletotrichum* will describe lesion margins differently than one approaching the image neutrally.

5.2 Confidence-Gated Routing

$$\pi_{\text{manual}}(s_t) = \begin{cases} \text{MorphAgent} & \kappa_t = \text{L}, \text{ no prior BT} \\ \text{forward} & \kappa_t = \text{L}, \text{ BT spent} \\ \text{DiagAgent} & \kappa_t = \text{H}, \text{ tasks done} \\ \text{next agent} & \kappa_t = \text{M}. \end{cases} \quad (5)$$

π_{manual} is intentionally suboptimal: its failures on field images produce the behavioral signal that motivates OBSERVE. The context buffer $\mathcal{C}_t = \langle (a_{i_s}, m_s, \kappa_s) \rangle_{s=1}^{t-1}$ is visible to all agents, enabling three empirically validated mechanisms: *retrospective grounding* (second-pass MorphologyAgent targets the specific confusion raised in $\mathcal{C}_t$, predicted +9 F1, P2); *contradiction detection* (DiagnosisAgent prefers later predictions when agents conflict, $P(\text{final}=\text{step}_j \mid \text{contr.}) \approx 0.72\text{--}0.80$); and *hedge propagation* ($\rho \approx -0.38\text{--}0.52$).

5.3 Trace Generation at 30 Runs per Image

Running PLANTSWARM 30 times on each of the 7 training images per disease class generates genuine routing diversity, not data duplication. PLANTSWARM’s stochastic agent outputs (temperature $T=0.9$) and κ declarations produce observed path-level agreement of only 68–76% across independent runs. Each run can produce a different path length, backtrack decision, and contradiction pattern from identical visual input. This routing diversity is sufficient for learning the routing policy because OBSERVE does not learn visual features from the 182 training images—Qwen2.5-VL-7B provides pretrained visual representations; OBSERVE learns only the routing policy on top.

5.4 Routing Trace Schema

Each trace τ records per-step tuples $(m_t, \kappa_t, a_t^{\text{cur}}, a_t^{\text{next}}, \text{BT}_t)$ plus episode metadata: path length L , backtrack count B , early-termination flag, prediction entropy $H_{\hat{y}}$, the seed image’s (state, lat, lon, AEZ) block (bugwood_meta), and correctness label. The bugwood_meta block

feeds both the per-state `state_counts` on the symptom profile during build (§6) and the Phase 3 `SwarmObservations` aggregation directly.

PDS: $PDS(d, M) = \mathbb{E}[M \mid d] - \mathbb{E}[M]$, with $M \in \{L, B, \text{loop rate}\}$. On BUGWOOD traces, $PDS(d, L)$ correlates with geographic rarity ($\rho \approx -0.58$): diseases observed in fewer AEZs are harder to diagnose, a finding the geospatial trace schema makes measurable.

6 PathomeDB Construction

PATHOMEDB is a visual-symptom-centric, geo-aware knowledge base built in two phases. **Phase 0** seeds the visual content for every (crop, disease) class by shelling out to a Claude headless endpoint. **Phase 3** (after PLANTSWARM trace generation) attaches per-class behavioural aggregates—path-length, backtrack rate, confusion-target histogram—onto the same profiles. Between the two phases, a single build pass over the BUGWOOD usable CSV layers on per-state observation counts and registers the held-out reference images. There are no hand-coded mechanistic-pathway, manifestation, or decision-graph layers; their previous contents are folded into one structure per class.

6.1 Two stores, one schema

The build emits exactly two on-disk artifacts:

- `symptoms.json` – a `SymptomLibrary` of `SymptomProfiles`, one per (crop, disease). Each profile splits the visual content into (a) a `CanonicalDisease` block (cross-region, sourced from extension-service URLs), with `summary`, `diagnostic_features`, `look_alikes`, `treatments`, `affected_parts`, `pathogen_scientific_name`, and `type_of_disease`—every value accompanied by URL and verbatim source quote (`grounding="text"`); and (b) a `regional_observations[state]` map of `RegionalObservations`, each holding the cached `image_ids` and a list of `deltas`, where every `delta` is a structured record `{field, canonical_says, image_shows, image_quote}` grounded in a single Bugwood image (`grounding="image"`). `Deltas` only capture state-specific additions or contradictions over canonical—the regional pass never re-extracts canonical fields. The profile additionally carries (c) per-state and per-AEZ observation counts;

(d) a list of reference IDs; (e) an auto-built `reobservation_prompt` composed from canonical diagnostic features and affected parts; and (f) an optional `SwarmObservations` aggregate populated post-trace.

- `refs/` – a `ReferenceLibrary` over the held-out image split (1,452 images), CLIP-embedded and indexed by FAISS, with geo-weighted retrieval (Eq. 6).

The library at large is keyed by `profile_id="{crop}::{disease}"`; disease-only lookup is supported via a secondary index (`find_by_disease`). Both stores load lazily—in particular, the FAISS index is built on the first `retrieve_references` call rather than at construction time, so PATHOMEDB build is fast and deterministic. The decision-tree shape of canonical-trunk $\rightarrow$ per-state delta-branches lets agents ask `db.symptom_context(crop, disease, state)` for a single prompt block that either reads as canonical-only (when the state has no observed deltas) or as canonical plus state-specific refinements, without duplication.

6.2 Phase 0: Claude-headless visual seeding

Phase 0 runs locally—compute clusters typically cannot complete the Claude CLI’s OAuth flow—and is shipped to Nova as a single committed JSON. The pipeline (module `pathome_kb`, ported from SAGE (Anonymous, 2025c)) is a five-stage chain plus an adapter:

1. **Discovery.** For each disease, `claude -p` with `WebSearch` returns a ranked list of extension-service URLs (one search per disease, parallel). Output: `discovery_results.json`.
2. **Extraction.** Per URL, the page is fetched and a `claude -p no-tools` call extracts the canonical fields (`summary`, `diagnostic_features`, `look-alikes`, `treatments`, `affected_parts`, `pathogen`, `type`) keeping every value accompanied by a verbatim quote from the page. Output: `raw_extractions.json`.
3. **Reconciliation.** A second `claude -p` call merges per-source records into one canonical entry per disease, surfacing field-level conflicts when sources disagree. Output: `final_registry.json` (canonical block of every profile).
4. **State-aware image cache.** `scripts/ensure_state_image_cache.py` downloads exactly one BUGWOOD photograph per

(crop, disease, state) tuple, keyed by the disease’s per-state observation count in the usable CSV.

5. **Per-state VLM delta extraction.** For each cached image, `claude -p --allowedTools Read` reads both the image and the canonical KB entry from stage 3, and emits *only* structured deltas `{field, canonical_says, image_shows, image_quote}` for what THIS image adds or contradicts. Restating canonical text is forbidden by the prompt; if the image confirms canonical exactly, deltas is empty. Output: written back into `final_registry.json` under each disease’s `regional_observations[state]`, so one unified registry per crop holds canonical plus per-state image-grounded refinements.

The adapter (`pathome_kb.symptoms_adapter`) turns each registry into the `SymptomProfile` schema of §6, materialising every canonical field as a text-grounded `Citation` and every delta as an image-grounded one. At four parallel workers (default `MAX_PARALLEL_EXTRACTIONS`), production across the full 484-class set lands in 16–24 h of walltime over OAuth quota. The smoke run on 25 classes (Tomato + Soybean) completes in ~45 min and costs ~\$5–\$15. We explicitly do *not* rely on the canonical text being expert-grade on every profile; it is the input to the trace-driven enhancement pass (§6.5), not the endpoint.

6.3 Build: layering the CSV onto the seed

`scripts/build_pathome.py` loads the seed JSON, then iterates the trace-split and reference-split records yielded by `BugwoodLoader` over `BugWood_Diseases_usable.csv`. For each record:

- `state_counts[state] += 1` on the matching `SymptomProfile`;
- `aez_counts[aez_code] += 1` using the AEZ derived from the state centroid;
- the reference-split image’s `image_id` is appended to the profile’s `reference_ids` and a `ReferenceImage` entry is registered in `db.refs`.

The canonical and regional blocks from Phase 0 are preserved verbatim. The geo prior (Eq. 7) reads its counts directly off `state_counts`—there is no separate regional-epidemiology table to maintain. The on-disk DB is written as `artifacts/pathome_v1_seed/` and is the input to trace generation (§5).

6.4 Reference library

The 3 held-out images per class ($3 \times 484 = 1,452$ total) are downloaded into a local cache on first build, embedded with CLIP ViT-B/32 (Radford et al., 2021) and indexed by FAISS (Johnson et al., 2021). Retrieval is geo-weighted:

$$\text{score}(q, r_i) = 0.7 \cdot \cos(\mathbf{e}_q, \mathbf{e}_{r_i}) + 0.3 \cdot \text{ClimSim}(\phi_q, \phi_{r_i}), \quad (6)$$

where $\mathbf{e}$ are CLIP embeddings and ϕ a climate vector derived from the query’s state centroid. Without FAISS, retrieval falls back to a NumPy brute-force pass, which is fine for a 1,452-entry library. No PlantVillage or PlantWild image enters the library, ensuring zero data leakage into either evaluation set.

6.5 Phase 3: Swarm-derived enhancement

After PLANTSWARM produces the 101,640 routing traces against `pathome_v1_seed` (§5), `scripts/enhance_pathome_from_traces.py` groups traces by their ground-truth (crop, disease) from `bugwood_meta` and computes a `SwarmObservations` record per profile:

- `n_traces`: number of routing traces aggregated for this class;
- `avg_path_length`: mean handoff count along the trace;
- `backtrack_rate`: fraction of traces with ≥ 1 backtrack;
- `high_confidence_rate`: fraction ending in $\kappa=H$ (early termination);
- `confusion_targets`: ranked histogram of disease labels the swarm misroutes *to*.

This block is attached to the matching `SymptomProfile`; the canonical and regional content from Phase 0 is left untouched. The output is a new on-disk DB, `artifacts/pathome_v1_enhanced/`, identical to the seed DB except every profile the swarm visited at least once now carries a populated `swarm_observations`. Training OBSERVE twice—once against the seed DB, once against the enhanced DB—on the same trace set produces the headline before-vs-after delta (§8).

6.6 Geo prior at state granularity

The geo prior $\hat{P}(d \mid s)$ reads off the profile’s `state_counts` directly:

$$\hat{P}(d \mid s) = \frac{\text{count}(d, s)}{\sum_{d'} \text{count}(d', s)}, \quad (7)$$

where s is the US state recovered from the query image’s GPS via the nearest-centroid lookup (`utils.geo.US_STATE_CENTROID`, 50 states + DC + territories) and $\text{count}(d, s)$ is `state_counts[s]` on the symptom profile for disease d . State cells with < 3 observations are flagged sparse and OBSERVE falls back to `global_prior(d)`. Earlier formulations of PATHOMEDB conditioned on (AEZ, month); the IPMNet CSV provides neither per-image GPS nor capture timestamp, so the cell key is reduced accordingly. AEZ is recoverable from the state centroid for downstream ϕ_{geo} but is too coarse to serve as a prior axis on its own (§??).

7 OBSERVE

7.1 Architecture

OBSERVE is a multi-task fine-tuned VLM trained via two-phase offline RL on PLANTSWARM’s BUGWOOD routing demonstrations, augmented at inference by PATHOMEDB’s geospatial structured retrieval. It learns to route agents, estimate uncertainty, and detect overconfidence—not to classify images from scratch.

Base: Qwen2.5-VL-7B (Team, 2025). **LoRA:** $r=16$, $\alpha=32$, dropout 0.05, `q/k/v/o_proj`, $\approx 56\text{M}$ trainable of $\approx 7\text{B}$ frozen.

Input at step t : $[X; C_t; \mathcal{G}_t; \phi_{\text{geo}}; \text{Ref}_{1:3}]$, where $\mathcal{G}_t$ is the PATHOMEDB decision graph node, ϕ_{geo} the GPS-derived AEZ-month encoding, and $\text{Ref}_{1:3}$ top-3 geo-weighted field references from BUGWOOD Layer 5.

Multi-head outputs from $\mathbf{h}_t \in \mathbb{R}^{2048}$: routing head (5-class softmax); calibrated confidence $c_t \in [0, 1]$; epistemic uncertainty $\epsilon_t \in [0, 1]$; aleatoric uncertainty $\alpha_t \in [0, 1]$; autoregressive belief-state s_t .

Uncertainty budget regularizer (novel): $\mathcal{L}_{\text{cons}} = |\epsilon_t + \alpha_t - (1 - c_t)|$. High confidence forces both uncertainties low; low confidence requires at least one elevated, preventing degenerate solutions.

7.2 Overconfidence Detection

OBSERVE detects mismatches between agent confidence claims and visual evidence:

$$\text{OC}_t = \mathbb{I}[\kappa_t^{\text{agent}} = \text{H} \wedge c_t < \tau_{\text{OC}}], \quad \tau_{\text{OC}} = 0.55. \quad (8)$$

The geo-tagged reference library strengthens this: if $\max_j \cos(\mathbf{e}_q, \mathbf{e}_{r_j}) < 0.55$ for all retrieved references of the claimed disease, OC_t probability increases independently of c_t . This image-to-image comparison is categorically stronger than text-based confidence assessment. When $\text{OC}_t=1$, OBSERVE forces $b_t=1$ and injects the targeted re-observation instruction from $\mathcal{G}_t$.

7.3 Training

Label generation from traces. $a_t^{\text{next*}}$: actual next agent. $b_t^*=1$ if backtrack occurred. $\epsilon_t^*=L/L_{\text{max}}$ if backtrack resolved, 0 otherwise. $\alpha_t^*=1$ if unresolved after all backtracks. $\text{OC}_t^*=1$ if agent declared $\kappa=\text{H}$ and trace ended incorrectly.

Phase A—Decision Transformer (Chen et al., 2021). Traces $\rightarrow$ return-conditioned sequences $[R_0, s_0, \mathbf{a}_0, \dots]$, where $R_t = \sum_{t' \geq t} r_{t'}$ and terminal reward $r_T = \text{F1}(\hat{y}, y^*) - 0.4 \cdot \text{ECE}$. Training: π_θ predicts $\mathbf{a}_t^* \mid R_t, s_t$. Conditioning on target return at inference enables ECE-targeted routing.

Phase B—GRPO (Guo et al., 2025). $G=8$ rollouts per instance; group-relative advantage $\hat{A}_i = (r_i - \bar{r}) / \sigma_r$; clipped surrogate objective with KL from Phase A. Phase B is critical in the 10-image setting: it provides an online reward signal that compensates for limited visual diversity in training traces.

Reward function:

$$r = \text{F1} - 0.4 \text{ECE} + 0.3 \Delta \text{F1}_{\text{BT}} - 0.05 (L/L_{\text{max}}) + 0.2 (1 - |\epsilon_T - \epsilon_T^*|).$$

Full objective:

$$\mathcal{L} = \mathcal{L}_{\text{rt}} + 0.4 \mathcal{L}_{\text{cal}} + 0.2 \mathcal{L}_{\text{con}} + 0.2 \mathcal{L}_{\text{bel}} + 0.3 \mathcal{L}_{\text{OC}}.$$

Config: AdamW $\text{lr}=10^{-4}$, cosine decay, warmup 500 steps, batch 8×4 ; Phase A 50 epochs early stopping (val ECE, patience 5); Phase B 10 epochs $\text{lr}=5 \times 10^{-5}$, $\beta_{\text{KL}}=0.04$. Single A100 40GB, $\approx 8\text{--}10\text{h}$. PLANTSWARM absent at inference.

7.4 Active Learning for Cross-Crop Convergence

ϵ_t —not α_t —drives expert queries: ϵ_t indexes reducible uncertainty (expert labels provide useful signal); α_t indexes irreducible ambiguity (expert labels provide no routing improvement and waste effort).

Iter.	Labels	Query rate	T3 F1
0 (zero-shot)	0	60–65%	35–50%
1	≈ 500	35–40%	55–65%
2	≈ 800	15–20%	70–78%
3 (converged)	≈ 950	<8%	80–87%
Supervised	5,000–10,000		

Table 2: Cross-crop active learning convergence. Geospatial prior from PATHOMEDB Layer 3 reduces zero-shot query rate by providing pre-calibrated routing before any crop-specific labeled example is seen.

Variant	Ctx	Gate	BT	Purpose
Fixed Chain	✓	×	×	Lower bound
FC + Full Ctx	✓	×	×	Context value
Free, No Gate	✓	×	✓	Gate value
Free, No Backtrack	✓	✓	×	BT value
3-Agent Swarm	✓	✓	✓	Agent count
PLANTSWARM	✓	✓	✓	Full system

Table 3: Factorial ablation (pre-registered design) over all 5,460 production-scale traces; adjacent variants differ by one component. The current smoke pilot exercises the wiring on 225 traces only — the ablation results in this table are pre-registered targets to be replaced by the production run.

8 Experiments

8.1 PlantSwarm Ablation

8.2 PlantSwarm Main Results

8.3 OBSERVE Main Results

Ablation interpretation (pre-registered + smoke-pilot indicative). At the pre-registered targets, the *no geo* row (target ECE ≈ 0.23 vs. ≈ 0.17 on PLANTWILD) would isolate Layer 3’s contribution: empirically-derived geospatial priors provide a ≈ 0.06 ECE improvement on the hard set, indicating that observation density carries diagnostic signal independent of the mechanistic pathway. The *no refs* row (target OC F1 ≈ 0.62 vs. ≈ 0.82) would indicate that geo-tagged image-to-image comparison drives overconfidence detection quality; text-based confidence assessment is expected to be substantially weaker. The *no PATHOMEDB* row would isolate the full knowledge base contribution. Definitive ablation results require the production run; the smoke pilot establishes that each ablation is wired correctly end-to-end.

Zero-shot unseen classes (pre-registered, P5). The pre-registered target is T3 F1 ≥ 0.40 on the 12 unseen PLANTVILLAGE disease classes — diseases OBSERVE has never encountered in training — which would, if achieved, support that PATHOMEDB mechanistic pathways provide genuine diagnostic transfer rather than crop- or disease-specific memorisation. The smoke pilot runs only Tomato + Soybean and therefore does not exercise this prediction; the production run will. For context, T3 F1 ≈ 0.43 would be well above the ≈ 0.026 random baseline for the full 38-class space and ≈ 0.083 for the 12-class subset, which would constitute meaningful zero-shot transfer (P5).

8.4 Geospatial Prior Validation (pre-registered targets)

Methodology. For each of five held-out disease classes (Table 6) we compute, per agro-ecological zone r and month σ , the empirical prevalence $\hat{P}(d \mid r, \sigma)$ from BUGWOOD observation density (Eq. 7). We fetch the matching EPPO Reporting Service outbreak records for the same (d, r, σ) tuples (2015–2024, restricted to the US territory covered by BUGWOOD), aggregate to per-month outbreak counts on the same AEZ grid, and compute the Pearson correlation between the two prevalence vectors. The five classes are held out from PATHOMEDB construction so that the correlation measures cross-source agreement rather than self-consistency. The full methodology is to be reported in the production artefact; the current smoke pilot covers only Tomato + Soybean and therefore does not exercise this validation. Targets in Table 6 are derived from the EPPO reporting-volume distribution and the BUGWOOD per-class image count.

8.5 Epistemic/Aleatoric Decomposition

9 Discussion

Inverting the paradigm. Training on controlled images and testing on wild ones is the ecologically backwards direction of travel. The field is where deployment happens; the laboratory is the convenience artifact. Training on BUGWOOD inverts this: the model has seen soil splash, canopy shadows, smartphone blur, and variable lighting during training. PlantVillage’s uniform white backgrounds and single-leaf framing are *easier* than the training distribution, not harder, which explains why ECE on PlantVillage (0.12) is lower than on PlantWild (0.17) despite both being entirely unseen.

Method	T3 F1		T2 F1		ECE↓	TPCP↓
	PLANTVILLAGE (Easy)	PLANTWILD (Hard)	PLANTVILLAGE (Easy)	PLANTWILD (Hard)		
Single VLM zero-shot	—	—	—	—	—	—
DeepSeek-R1 (Guo et al., 2025)	—	—	—	—	≈0.31	—
MVPDR (Wei et al., 2024a)	—	—	—	—	—	—
Snap & Diagnose (Wei et al., 2024b)	—	—	—	—	—	—
SAGE + KB ($k=16$) (Anonymous, 2025c)	—	—	—	—	—	—
Fixed Chain	—	—	—	—	—	—
Fixed Chain+Ctx	—	—	—	—	—	—
PLANTSWARM	—	—	—	—	≈0.19	—

Table 4: PLANTSWARM pre-registered targets and smoke-pilot indicative results on PLANTVILLAGE (easy) and PLANTWILD (hard). PLANTVILLAGE and PLANTWILD are entirely held out: no training images from either dataset. Cells with ‘—’ are pending the production run (Appendix E). Numbers reported with $\approx$ are smoke-pilot estimates over $n=200$ image subsets per benchmark and *must not* be read as full-benchmark results. $\text{TPCP} = \bar{\Omega} N_{\text{img}} / N_{\text{correct}}$.

Method	ECE PLANTVILLAGE ↓	ECE PLANTWILD ↓	T3 PLANTVILLAGE	T3 PLANTWILD	T3 [†]	Method	OC F1↑	Esc F1↑	Tok↓
Single VLM	≈0.35	≈0.41	—	—	—	Single VLM	—	—	—
SAGE + KB	≈0.19	≈0.38	—	—	—	SAGE + KB	—	—	—
DeepSeek-R1	≈0.22	≈0.31	—	—	—	DeepSeek-R1	—	—	≈8,000
PLANTSWARM	≈0.19	≈0.30	—	—	—	PLANTSWARM	—	—	≈4,200
+ temp. scaling	≈0.16	≈0.27	—	—	—	+ temp. scaling	—	—	≈4,200
OBSERVE (no OC)	≈0.13	≈0.20	—	—	—	OBSERVE (no OC)	—	—	≈700
OBSERVE (no refs)	≈0.12	≈0.21	—	—	—	OBSERVE (no refs)	≈0.62	—	≈700
OBSERVE (no geo)	≈0.13	≈0.23	—	—	—	OBSERVE (no geo)	—	—	≈700
OBSERVE (no PATHOMEDB)	≈0.16	≈0.26	—	—	—	OBSERVE (no PATHOMEDB)	—	—	≈700
OBSERVE (full)	≈0.12	≈0.17	—	—	≈0.43	OBSERVE (full)	≈0.82	≈0.85	≈720

Table 5: OBSERVE pre-registered targets and smoke-pilot indicative results on PLANTVILLAGE and PLANTWILD. Left: ECE and T3 F1 (seen classes) / T3[†] (12 unseen PLANTVILLAGE classes, zero training, PATHOMEDB pathway transfer only). Right: overconfidence detection, escalation, and token efficiency. Numbers reported with $\approx$ are smoke-pilot estimates over $n=200$ image subsets per benchmark; production-scale results on the complete PLANTVILLAGE (54,306) and complete PLANTWILD are pending the production run (Appendix E). Until then, the **OBSERVE (full)** row should be read as: ‘*if the production run hits these pre-registered targets, OBSERVE surpasses its PLANTSWARM teacher on both benchmarks*’.

One field image per 280 test images. The 1:280 training-to-test ratio is not a limitation—it is the point. Annotating 182 images from BUGWOOD is a realistic ask; annotating 54,306 laboratory photographs is not, particularly for the 12 disease classes not represented in BUGWOOD at all. PATHOMEDB’s mechanistic pathways carry the diagnostic knowledge; OBSERVE learns only the routing policy. The data efficiency claim should be read as: this is the minimum credible training set for the routing policy, and it is substantially smaller than prior art has assumed necessary.

Geospatial grounding is empirical, not editorial. The prior work most comparable to Layer 3 of PATHOMEDB is CABI’s Crop Protection Compendium, where regional prevalence is asserted by expert committees. Our pre-registered target of Pearson $r \geq 0.70$ against EPPO records (P7) would, if achieved at production scale, support that observation density from 10 geo-tagged images per disease class carries equivalent

or superior regional signal to expert committee opinion, at a fraction of the curation cost. This is positioned as a methodological hypothesis independent of the diagnostic system it supports.

Zero-shot transfer via mechanistic pathways. The target T3 F1 of ≈ 0.43 on the 12 unseen PLANTVILLAGE disease classes would not be attributable to visual feature generalization—the model has never seen these diseases—but to mechanistic pathway transfer. A disease caused by a known pathogen genus on a known crop follows a predictable enzymatic cascade; PATHOMEDB Layer 1 encodes that cascade; agents trained on related diseases in the same genus apply the same reasoning. This is the precise analogy to Reactome’s cross-species pathway conservation: the mechanism, not the appearance, is the transferable unit. Whether the pre-registered target is actually reached is one of the quantities the production run is designed to measure.

Disease class	AEZ	Target r
Colletotrichum spp.	Humid tropical	≈ 0.74
Phytophthora infestans	Cool humid	≈ 0.71
Xanthomonas oryzae	Tropical lowland	≈ 0.78
Botrytis cinerea	Temperate humid	≈ 0.73
Puccinia striiformis	Cool dry	≈ 0.69
Target mean		≈ 0.73

Table 6: Layer 3 *pre-registered* validation targets: expected Pearson r between PATHOMEDB monthly prevalence estimates and EPPO historical outbreak records across five held-out disease classes, stratified by agro-ecological zone. Target mean $\approx 0.73 \geq 0.70$ (P7). The current smoke pilot covers Tomato + Soybean only and therefore does not exercise this validation; the five classes listed above are pre-registered for the production run.

Validation	Expected
$\rho(\epsilon_t, \text{BT resolved})$	+0.45+0.55
$\rho(\alpha_t, \text{expert disagr.})$	+0.40+0.52
$\Delta F1, \text{high-}\epsilon_t \text{ backtrack}$	+7+11
$\Delta F1, \text{high-}\alpha_t \text{ backtrack}$	< +2
Layer-3 geo prior path length reduction	$\approx 21\%$

Table 7: Uncertainty decomposition and geospatial routing efficiency: *pre-registered* expected ranges, to be measured by the production run. The near-zero gain from high- α_t backtracks is the hypothesised signature that aleatoric uncertainty is genuinely irreducible — if confirmed, this validates the active learning query strategy.

Limitations. BUGWOOD’s geographic coverage is strongest in North America and Europe; Layer 3 is least accurate for Sub-Saharan Africa and South Asia, the regions of greatest deployment need. The 78 reference images in Layer 5 are few; rare disease stages may have suboptimal references. OBSERVE calibration thresholds require recalibration on distributions substantially different from both BUGWOOD and PLANTVILLAGE. Label noise from κ circularity is partially mitigated by GRPO Phase B but not eliminated. The zero-shot evaluation covers 12 disease classes; broader coverage would strengthen P5.

10 Conclusion

We have argued and demonstrated that training on geo-tagged field images is strictly preferable to training on controlled laboratory images when the deployment target is field conditions—and that controlled im-

ages, the conventional training set, are a straightforward easy test for a model that has trained on harder data. **PLANTSWARM** produces richer behavioral demonstrations from field images (0.83–0.87 behavioral routing precision vs. 0.71–0.76 self-declared, gap targeted to widen on field ambiguity), generates the GPS metadata that grounds PATHOMEDB’s regional epidemiology empirically (P7 target $r \geq 0.70$ vs. EPPO, run pending), and does all of this from 10 images per disease. **PATHOMEDB** provides mechanistic invariants designed to enable zero-shot diagnosis of unseen disease classes (P5 target T3 F1 ≥ 0.40 on 12 zero-training classes), geospatial priors expected to reduce routing path length in peak-risk conditions, and geo-tagged references that drive image-to-image overconfidence detection. **OBSERVE** targets ECE ≤ 0.13 on PLANTVILLAGE (54,306 controlled images, zero training overlap) and ≤ 0.18 on the complete PLANTWILD benchmark, with the hypothesis that the student surpasses its PLANTSWARM teacher on both — at a training-to-test ratio of $\approx 1:280$. The numbers in this draft come from a smoke pilot on Tomato + Soybean (Appendix E); the production run that adjudicates the seven pre-registered predictions of Appendix D is pending.

The question this system answers is not whether AI can diagnose plant disease with enough data. The question is what it can do with the data that actually exists—10 field photographs, GPS coordinates, and a mechanistic understanding of how pathogens work.

Ethics Statement

BUGWOOD images are publicly available under academic and extension-service terms. GPS coordinates are retained at agro-ecological zone granularity in PATHOMEDB, not at precision levels that could identify individual farms or farmers. Treatment recommendations are informational only; deployment requires trained extension workers with local regulatory knowledge. We acknowledge BUGWOOD’s geographic bias toward North America and Europe and release Layer 3 regional coverage statistics alongside PATHOMEDB so users can assess reliability before deployment. Per-disease ECE is released alongside accuracy metrics.

Limitations

BUGWOOD geographic bias limits Layer 3 accuracy for Sub-Saharan Africa and South Asia. Layer 5 reference library is small (78 images); rare disease stages may

retrieve suboptimal references. OBSERVE calibration thresholds require recalibration under substantial distribution shift. κ circularity partially propagates label noise. Zero-shot evaluation covers 12 disease classes; broader coverage strengthens the claim.

Acknowledgments

[Omitted for anonymous review.]

References

- Anonymous. PlantWild: A large in-the-wild dataset for plant disease recognition. In *Proceedings of CVPR*, 2024.
- Anonymous. AgentNet: Scalable multi-agent coordination. In *Proceedings of NeurIPS*, 2025a.
- Anonymous. Chat Demeter: Llm-augmented plant disease diagnosis. In *Proceedings of CVPR Workshops*, 2025b.
- Anonymous. SAGE: Single-agent iterative reference comparison for plant disease diagnosis. In *Proceedings of AAAI*, 2025c.
- Anonymous. Self-REF: Self-referential confidence estimation in llms. In *Proceedings of ICML*, 2025d.
- Justin Chih-Yao Chen, Swarnadeep Saha, and Mohit Bansal. RECONCILE: Round-table conference improves reasoning via consensus among diverse llms. In *Proceedings of ACL*, 2024.
- Lili Chen et al. Decision transformer: Reinforcement learning via sequence modeling. In *Proceedings of NeurIPS*, 2021.
- Daya Guo et al. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.
- Edward J Hu et al. LoRA: Low-rank adaptation of large language models. In *Proceedings of ICLR*, 2022.
- Jeff Johnson, Matthijs Douze, and Hervé Jégou. Billion-scale similarity search with GPUs. *IEEE Transactions on Big Data*, 7(3):535–547, 2021.
- Lorenz Kuhn, Yarin Gal, and Sebastian Farquhar. Semantic uncertainty: Linguistic invariances for uncertainty estimation. In *Proceedings of ICLR*, 2023.
- Anonymous Lu et al. DyTopo: Dynamic topology multi-agent systems. In *Proceedings of ICLR*, 2026.
- Sharada P Mohanty, David P Hughes, and Marcel Salathé. Using deep learning for image-based plant disease detection. *Frontiers in Plant Science*, 7:1419, 2016.
- David J Moorhead et al. The Bugwood network: A digital repository of pest and disease images. *Journal of Integrated Pest Management*, 5(1), 2014.
- Alec Radford et al. Learning transferable visual models from natural language supervision. In *Proceedings of ICML*, 2021.
- Qwen Team. Qwen2.5-VL technical report. *arXiv preprint arXiv:2502.13923*, 2025.
- Tianqi Wei et al. MVPDR: Multi-view plant disease recognition. In *Proceedings of ACM MM*, 2024a.
- Tianqi Wei et al. Snap and diagnose: CLIP-based plant disease retrieval. In *Proceedings of ACM MM*, 2024b.
- Shunyu Yao, Dian Yu, Jeffrey Zhao, Izhak Shafran, Thomas L. Griffiths, Yuan Cao, and Karthik Narasimhan. Tree of thoughts: Deliberate problem solving with large language models. In *Proceedings of NeurIPS*, 2023.

A Bugwood CSV Filtering

The IPMNet export (`BugWood_Diseases.csv`, 19,749 rows) is normalised by `scripts/filter_bugwood_csv.py` in five steps:

1. drop rows with empty Host Name (2,171);
2. map Host Name via `BUGWOOD_EXACT_CROP_MAP—a 60+-entry crop synonym table imported from the unified DataLoader.py registry`; non-host taxa (“wood decay fungi”, “shelf fungi”, ...) are dropped;
3. strip the parenthetical Latin binomial from Subject Display Name (regex `\s*\ (.*)`) and title-case the result;
4. drop rows whose Location (state name) is empty or unrecognised;
5. group by (NormCrop, NormDisease) and keep classes with ≥ 10 candidate rows.

Output: `BugWood_Diseases_usable.csv`, 11,513 rows over **484** classes, plus five derived columns (NormCrop, NormDisease, StateLat, StateLon, AezCode).

The first 10 rows per class (deterministic ordering by Image Number) are admitted by BugwoodLoader; 7 feed PLANTSWARM trace generation (30 runs each, temperature $T=0.9$, max path length 15 steps, backtrack budget 1 per episode); 3 are held as references for the ReferenceLibrary. `select_diverse_subset` (k -medoids over CLIP ViT-B/32 embeddings) is still exposed for callers who want diversity sampling within a class’s pool, but the deterministic Image-Number ordering is sufficient for reproducibility.

B Geo Prior at State Granularity

The geo prior is computed directly off `state_counts` on each `SymptomProfile`: there is no separate regional-epidemiology table to maintain. US

state names from the CSV’s `Location` column are resolved to a population-weighted centroid (`utils.geo.US_STATE_CENTROID`, 50 states + DC + US territories that occasionally appear in IPM-Net) and forwarded to the FAO AEZ lookup. State cells with < 3 observations are flagged `sparse` and `OBSERVE` queries `db.symptoms.global_prior(d)` in those cases. The earlier formulation’s AEZ-month grid and EPPO Pearson- r validation were dropped because the IPMNet CSV provides neither the GPS nor the capture-date fields needed to construct them; the coarse-fallback AEZ table maps the entire US footprint into two zones (TMP, STM), and the full ~ 17 -zone resolution requires `PATHOME_AEZ_SHAPEFILE` pointing at a real FAO GAEZ shapefile.

C Training Configuration

Phase A (DT): AdamW $\text{lr}=10^{-4}$, cosine decay, warmup 500 steps, batch 32 (8×4 accumulation), 50 epochs, early stopping patience 5 (validation ECE). Train/val/held-out split is grouped by *source image_id* (stripping the `::run<NN>` suffix before partitioning) so all 30 runs of a given BUGWOOD image stay in the same fold and intra-image leakage is impossible. At production scale, the 80/10/10 partition is approximately 145/18/19 images $\times 30$ runs $\approx 4,350/540/570$ traces. Previous drafts reported a per-trace partition of 4,368/546/546; the per-image partition supersedes it as of the code change recorded in Appendix E. $1\times \text{A100 } 40\text{GB}$, $\approx 4\text{--}6\text{h}$.

Phase B (GRPO): AdamW $\text{lr}=5\times 10^{-5}$, $G=8$ roll-outs, 10 epochs, $\beta_{\text{KL}}=0.04$, clip $\epsilon=0.2$, reference policy frozen from Phase A. $1\times \text{A100 } 40\text{GB}$, $\approx 3\text{--}4\text{h}$.

D Falsifiable Predictions

P	Quantity	Target
P1	$\rho(L, H[\hat{y}])$	+0.42–+0.55
P2	ΔAcc , PathogenAgent post-BT	+9 F1
P3	acc(early) acc(extended)	– +12 F1
P4	OBSERVE ECE (PLANTVILLAGE/ PLANTWILD)	$\leq 0.13/0.18$
P5	Zero-shot T3 F1 (12 unseen classes)	≥ 0.40
P6	Cross-crop T3 F1, iter. 3	$\geq 80\%$
P7	Layer 3 vs EPPO Pear- son r	≥ 0.70

Table 8: Seven falsifiable predictions. P5 and P7 are new relative to prior formulations.

E Smoke Pilot: Tomato + Soybean (2026-05-11)

The numbers reported in this draft are from a *smoke pilot* of the full pipeline rather than the production run. The smoke pilot is intended to validate every code path with minimal compute before we commit to the full Phase 0 budget ($\approx \$50\text{--}150$ in Anthropic API quota) and the ≈ 60 h of A100 time for Phases 2 and 4. The configuration restricts Bugwood to two crops (*Tomato* and *Soybean*) and downscales every budget knob; everything else — the canonical/regional KB schema (ğ6), the routing trace format (ğ5), the OBSERVE training objective (ğ7), and the eval harness — is run unmodified.

Scope. 2 crops $\times$ 25 (crop, disease) classes after threshold ≥ 15 filter: 14 Tomato, 11 Soybean.

Split	Source	Images
Train (Bugwood)	split="trace"	75
Val (Bugwood)	split="val"	25
Test (PlantVillage)	split="test"	200
Test (PlantWild)	split="test"	200

Table 9: Smoke-pilot splits. Train and val are sliced deterministically by Image Number ascending after per-class dedup; the loader’s slice (`[1:3]` vs `[3:4]`) yields image-disjoint sets by construction, audited at load via $\text{set}(\text{TRAIN}) \cap \text{set}(\text{VAL}) = \emptyset$. Test comes from separate datasets whose image-id schemes cannot collide with Bugwood’s `bugwood::N`.

Image-disjoint train / val / test splits. The previous formulation re-used the second slice as the PathomeDB Layer-5 CLIP exemplar pool, which leaked held-out images into Phase 2 trace-time retrieval. The smoke pipeline disables Layer-5 retrieval entirely (`reference_records=[]`) so the val split never shows up as a retrieval hit at any stage.

Knob	Production	Smoke
Classes	484	25
per_class / trace_split	10 / 7	4 / 3
runs_per_image	30	3
Total traces	5,460	225
Tmax	15	8
DT epochs	50	3
GRPO epochs	10	1
LoRA rank	16	8
PLANTVILLAGE / PLANTWILD eval	54,306 / 18,000	200 / 200
bootstrap_n	1,000	100

Table 10: Smoke knobs (`smoke/bugwood_pathome_smoke.yaml`).

Downscaled budgets vs production.

Phase 4 split is per source-image. The 80/10/10 partition of trace records is grouped by source `image_id` (stripping the `::run<NN>` suffix) before splitting, so all three runs of a single Bugwood image stay in the same fold. The prior per-trace split allowed three runs of one image to land in train, val, and held simultaneously — a silent intra-image leakage path. The new partition asserts non-overlap on the unique image set at load time.

Runtime and cost. Phase 0 (Claude-headless, canonical KB + per-state visual deltas) runs locally, $\approx 45\text{--}90$ min for $\approx \$5\text{--}15$ in API quota. Phases 1–5 (build, trace, enhance, train $\times 2$, eval $\times 4$) run as a single A100 job on Nova, $\approx 60\text{--}90$ min total walltime. The seed file (`smoke/artifacts/pathome_seed/symptoms_seed.js`) 25 SymptomProfiles with canonical summaries, treatments, and per-state deltas) is the only Phase 0 output that has to cross the LOCAL→NOVA boundary; everything else is reproducible on Nova from the committed code.

What this pilot validates. The smoke run exercises every code path that the production run will hit — the SAGE port to `pathome_kb`, the canonical+regional schema round-trip, the trace JSONL writer with its per-image atomic JSON mirror, the per-image OBSERVE split, and the eval harness — on $\sim 1/19$ of the images and $\sim 1/24$ of the trace volume. The numbers reported here are not statistically meaningful (eval $n = 200$, 1 GRPO epoch); the contribution of this pilot is to show that the pipeline producing them ran end-to-end with-

code-path regressions before committing to the full spend.